

Little evidence of ideological bias in YouTube suspensions of political content during the 2024 US Election

Andreu Casas*

Summary

Social media play a crucial role in moderating speech. In the United States, conservative elites claim that platforms unfairly sanction conservative voices. Yet systematic, large-scale analysis is lacking. For more than a year leading up to the 2024 Election, we monitored 20,000 YouTube channels posting about US politics, and their millions of videos. By election day, the platform had removed 2.17% of the channels and 3.3% of the videos. Conservative channels and videos were suspended 3-4 times more frequently than liberal ones. However, conservative videos accounted for most the hateful content (58.18% v. 23.35%, for liberal videos) and misinformation and conspiracies (71.9% v. 10.08%), as identified using fine-tuned Vision Language Models. Accounting for platform violations explains most of the ideological gap in suspensions, leaving a residual difference of less than a percentage point, suggesting that ideological differences reflected asymmetries in compliance with platform policy far more than a uniform ideological moderation bias.

*Director of the Oxford Social Media Observatory, University of Oxford (Department of Politics and International Relations; Oxford Internet Institute; Center for Advanced Social Science Methods): andreu.casas@politics.ox.ac.uk.. I thank Freek Cool for the amazing support during the data collection process, and to Georgia Dagher, Irene Dekker, and Sanne Wiering for their amazing research assistance.

Main

In the last decades, social media companies have gained an unprecedented power to regulate speech through their content moderation policies (1; 2). Among these, YouTube stands out as one of the most used platforms in the United States (US), particularly for consuming news and politics (3). In recent years, conservative elites have argued that mainstream platforms unfairly moderate conservative voices, especially during the Biden Administration. Senator Ted Cruz stated that “the pattern of political bias from YouTube and Google is massive” (4), Representative Jim Jordan argued that “Big Tech is out to get conservatives” (5), and in an executive order, President Trump stated that “the previous administration trampled free speech rights by censoring Americans’ speech on online platforms” (6). According to a 2020 survey from PEW (7), around 90% of Republicans believe that social media platforms intentionally censor political views they find objectionable (v. 56% for Democrats). However, despite the urgent implications of these claims for US democracy and beyond, we still lack systematic analyses of the suspension of content and accounts on mainstream social media platforms.

The lack of research in this area primarily stems from the complexity of collecting and analyzing moderation multimodal data from platforms at scale. Suspensions happen but not in high numbers (8; 9), which means that large amounts of data must be collected to have enough statistical power. Suspended accounts and content become immediately inaccessible, requiring data to be collected as close to real time as possible. Suspensions arise for reasons beyond ideological leaning, meaning that a range of confounding factors must be taken into account (9). The latter is particularly challenging given the multimodal nature of the data (composed of text, images, videos, and sound), as some key moderation predictors (e.g. misinformation) can be encoded in the text but also in the visual of a post (10).

Here, we use a variety of state-of-the-art computational methods to address these chal-

lenges, and to critically assess suspension patterns among 20,000 YouTube channels and the 6 million videos they posted, for more than one year leading up to the 2024 US Election. Our main question is whether suspension patterns were reflective of an ideological bias. We use computational methods to estimate the ideological leaning of the accounts and videos, as well as to measure other account and video level confounders (such as account coordination, and whether videos contain hate speech or misinformation), and estimate statistical models predicting account and video suspension as a function of a wide range of theoretically relevant predictors.

We expand past literature in several ways. First, prior research on account and content suspension has often been limited in scale or topical scope, for example, focusing on a single geopolitical issue (9) or a single referendum debate (11). Here, we study suspensions among a large sample of accounts and videos addressing a wide range of (non)political topics, for a long period of time leading up to an important election, and on one of the platforms most often used for news and politics: YouTube (3). Second, while some recent large-scale work has begun to account for confounders beyond ideology when explaining suspension asymmetries (for example, using domain-level news quality ratings and off-the-shelf classifiers of bot-likeness and toxic language (8)), these approaches necessarily rely on proxies for the underlying behavior of interest, since assessing the actual content of individual posts at scale has not been feasible. Here, we measure these confounders directly: we fine-tune open-weights machine learning models to automatically identify, in the millions of videos we study, whether they discuss politics, their specific topic, and whether they contain hateful content, misinformation, or conspiracies. Third, existing research has mostly ignored the visual information increasingly embedded in social media messages, relying on text or metadata alone (8). This can be problematic given that some relevant predictors of suspensions, such as misinformation, can be encoded in the visual part of the post (10). Here, we use Vision Language Models (VLMs) as the open-weights models we fine-tune, allowing us to jointly leverage both the transcript


and visual frames of each video. Fourth, prior important research looks at suspension patterns in non-Western contexts, such as China (12), but less is known about account and content removal in Western democracies (13). Here, we provide detailed insights about the moderation of political content on a Western platform during a pivotal democratic election. Finally, the few studies that do look at Western countries, such as France, UK, and Germany (14), and the US (8), mainly focus on removals among ordinary accounts. Here, we pay particular attention to removals of salient content producers, who often dominate the narratives that circulate on social media networks.

Results

A large sample of politically-interested YouTube channels

Between July 1st 2023, and November 6th 2024, we monitored 20,343 YouTube channels that posted, at least on occasions, about US politics. Figure 1.A provides an overview of the snowballing technique used for constructing this sample of channels, and further details can be found in the *Methods* section. In a nutshell, the starting point is a set of seed elite channels (members of Congress, the U.S. administration, mainstream media, and journalists; $N = 289$). We then expanded the sample, and its ideological and stylistic diversity, by including other politically relevant channels to which users commenting on political videos from these elite channels were also subscribed. The final sample includes channels from political elites but also influencers, alternative media, entertainment figures, podcasters, history and educational channels, etc. In total, during the period of analysis, we collected and tracked 6,068,468 videos from 11,979 of the channels. The remaining channels in the sample did not post any videos during this time period. On average, we collected 12,284 videos a day (95% Confidence Interval: 11,966-12,602), for an average of 508 videos per active channel for the whole time

Figure 1: **A.** Flowchart describing the snowball method we used to build the sample of channels analyzed in the paper **B.** Percentage of the 20,343 analyzed channels removed by YouTube between August 2023 and November 2024 (by ideology of the channel). **C.** Cumulative number of suspended channels by ideology, from August 2023 through November 2024. **D.** Percentage of channels suspended each month, by ideology.



period (95% CI: 494-522).

Our interest here is to use the channels in this sample (and the videos they posted) to explore potential ideological differences in channel and video suspensions. As detailed in the *Methods*, for robustness, we use two different measures of ideology. The first one is a network-based measure at the channel level. We follow an extensive literature in the study

of social media and politics (15; 16), and build on the assumption that users with a similar ideology “follow” similar accounts, to estimate the ideology of all channels into a left-right continuum. In some descriptive figures we divide the continuous scores into liberal, moderate, and conservatives channels, to more easily communicate the findings. The other is a content-based measure at the video level. We use the video transcripts and frames to fine-tune an open-source VLM with human-annotated data of our own, to classify YouTube videos into neutral, liberal, moderate, and conservative (see further details and performance in the *Methods* and Appendix).

Descriptive patterns of suspension and platform violations


During the period of analysis the platform removed 441 of the 20,343 monitored channels (2.17%). Figure 1.B breaks down the results by the ideological leaning of the channels, as estimated by the network-based measure described in the *Methods* (continuous score broken down into three categories here for interpretability). The suspension rate is the highest among conservative channels, 3.3%, compared to 2.36% for moderate and 1.07% for liberal channels. Note that the final sample had a similar number of channels in each of these ideological bins, around 6,000-7,000, with the largest number of channels in the liberal category ($N = 7,499$). Figure 1.C shows the cumulative number of suspended channels during the period of analysis, by ideology. Suspension rates remained rather constant across the whole time period, with no clear sudden increase/decrease at any particular point. Additionally, despite some small variations, Figure 1.D shows that the ideological distribution of the suspended channels also remained stable during this time period, with conservative channels always accounting for 40-50% of the monthly channel suspensions, whereas liberal channels rarely accounted for more than 25% of the channel suspensions in a given month. Figure 1.E breaks down the suspension rates by the topical focus of the channels. These are topical labels chosen by the

channel owners themselves, and they are not mutually exclusive: a channel can have more than one topic label and be present in more than one of the bars in the figure. The results show that other channel-level characteristics, beyond the ideological slant of the channel, can be predictive of suspension. The figure for example shows the highest suspension rates for military channels (14.86%: 66 out of 444 in the sample). These are channels showing footage from historical military conflicts (e.g. Vietnam, Afghanistan), as well as of military equipment (e.g. the latest acquisitions by the US military). These channels are probably suspended for portraying violent and sensitive content.

The platform also removed 210,245 of the 6,296,227 videos collected during the period of analysis (3.3%). Note that all videos from a suspended channel are also automatically suspended. None of the videos belong to channels in the original seed list of elite channels (politicians, news organizations, journalists). Figure 2 reports suspension patterns for English videos collected for a one year period: between November 2023 and November 2024 ($N = 2,325,867$). We zoom into these English videos because they are the ones for which we can generate more reliable estimates for the set of theoretical quantities of interest, such as whether a video is about politics, the particular issue area, and whether they spread misinformation or hateful content. The overall descriptive patterns in the top panels in Figure 2 (A,B,C) are very similar to the channel suspension patterns in Figure 1. Figure 2.A contains results for non-political videos whereas B and C report findings for political videos. The results in Figure 2.A and B are based on the network-based ideology scores at the channel level, and so on the assumption that the ideological slant for the channel can be propagated to all videos from that channel; whereas the numbers in Figure 2.C are based on the content-based ideology scores at the video level.

Three main findings stand out in these top panels. First, suspension rates are higher for political (2.88% in Figure 2.B and 2.66% in Figure 2.C) than non-political videos (2.11% in Figure 2.A). Second, conservative political videos are suspended at substantially higher rates

Figure 2: Suspension rates for *English* videos collected between November 2023 and November 2024. **A.** Non-political videos, by ideology of the channel (continuous score split into Liberal, Moderate, and Conservative). **B.** Political videos, by ideology of the channel. **C.** Political videos, by ideology of the video. **D.** All videos, by whether they contain hateful content or not; and by whether they spread misinformation or conspiracies, or not. **E.** All videos, by hateful/misinformation content, and by ideology of the video.



(4.72-4.92%) than moderate (2.47-2.87%) and liberal videos (1.12-1.49%). Finally, the previous findings are robust to using different measures of ideology, as the patterns look similar when using the network-based measure at the channel level, and the content-based measure at the video level.

However, the two bottom panels of Figure 2 (D,E) also show crucial ideological differences with respect to the toxicity of the content posted by these channels. As expected, Figure 2.D shows higher suspension rates for videos with hateful content (4.97% v. 2.43% for non-hateful) and misinformation (8.21% v. 2.87% for non-misinformation). Figure 2.E shows that conservative videos account for a small proportion of non-hateful (8.85%) and non-misinformation (17.35%) videos. However, in line with existing research (8), conservative videos account for most of the videos that spread hateful content (58.18%) and misinformation (71.9%). Arguably, these differences in the supply of toxic content, rather than an ideologically biased implementation of the platform’s terms of service, could then account for the higher suspension rates observed among conservative channels and videos.


A closer look at hateful and misinformation content

For a better understanding of the distribution of potential suspension confounders, Figures 3.A and B provide a more detailed picture of the videos with hateful content and misinformation and conspiracies, which violate YouTube’s terms of service.

In Figure 3.A videos are broken down by ideological deciles, with the same number of videos in each of the ten ideologically sorted bins. By further breaking down the continuous channel ideology score into a finer set of bins, we get a more nuanced picture of what is happening. Both for hateful content, and for misinformation/conspiracies, we observe an overall U-shape pattern, with videos in the middle moderate part of the distribution being the least toxic, and with toxicity increasing in both extremes of the distribution. However, both for hate speech and misinformation, we observe substantially higher toxicity levels in the most conservative decile. In line with prior findings and theorizing (17), the analysis points to a greater supply of toxic content on the extreme right that is not mirrored on the extreme left.

Figure 3.B shows the most predictive words for video transcripts annotated as hateful, and

Figure 3: **A.** Proportion of videos from channels from each ideological decile that are hateful (top) and spread misinformation (bottom). **B.** Most predictive words for hateful and misinformation/conspiracy videos.



as spreading misinformation and conspiracies. We use the output of more transparent bag-of-words text classifiers in this figure: two logistic regressions trained on the approximately 4,000 transcripts we have manually annotated (see Appendix B for a detailed description of the annotated data). For hateful, words around gender, and masculine/manosphere discourse appear at the top: e.g. transgender, feminine, pronouns, woman, girls, trans, etc.; a narrative currently dominated by the far right. For misinformation, we see words related to a variety of conspiracies. For example, “genes” is used in videos about some governments allegedly developing secret biological weapons to eradicate particular ethnicities, “jews”/“israelis” is used in videos spreading misinformation around the Israel-Palestine conflict, and “rigged” in videos spreading misinformation about the 2020 presidential election. These are also narratives more present among (extreme) conservative discourses.

Disentangling ideology from other predictors of suspension

The previous findings indicate that a variety of factors beyond ideology, such as the topical focus of the channel (Figure 1.E), or platform violations such the use of hate speech or the spread of misinformation (Figure 2.D), are also correlated with suspension. Here, we estimate a set of logistic regressions in order to assess whether ideology is a significant and substantive predictor of channel and video suspension once we account for other relevant confounders. For a more robust test, we focus on English videos posted between November 2023 and election day in November 2024 (~ 2.3 million, from ~ 8.5 k channels); the same data shown in Figure 2. We do so because although the multimodal VLM used for identifying some key video measures (e.g. hateful content and misinformation) is multilingual in nature, previous research shows that these large models are substantially more accurate in English than other languages (18).


In some of the models we predict channel suspension (Models 1 and 4) and in some others video suspension (Models 2, 3, 5, and 6), both binary outcomes. For a cleaner test, the models predicting video suspension only include videos from non-suspended channels (as for suspended channels, all videos are immediately suspended, and we do not clearly know which one/s motivated the channel suspension). In these models predicting video suspension, standard errors are clustered at the channel level, to account for the nested nature of the data. Additionally, some of these video models include all videos (Models 2 and 5) while some others only include those predicted by our VLM to be about US politics (Models 3 and 6). The most consequential difference is that some models only include channel and/or video ideology as predictors (Models 1, 2, and 3), whereas others also include a range of theoretically relevant covariates (Models 4, 5, and 6). As for these additional covariates, apart from the continuous network-based channel ideology score, the models predicting channel suspension include: number of days since the creation of the account (log), number of subscribers (log), number of videos ever posted (log), number of views ever received (log), mean and average coordina-

tion scores, the total number of videos collected for this study (log), as well the number that are political (log), hateful (log), and spread misinformation or conspiracies (log), and channel topic-fixed effects (23 topic categories). The models predicting video suspension, apart from including a categorical variable indicating the ideological leaning of the video (liberal, moderate, conservative, or neutral), they also include a set of channel-level covariates: the continuous network-based channel ideology, the type of channel (politician, media, or other), number of days since the creation of the account (log), number of subscribers (log), number of videos ever posted (log), number of views ever received (log), mean and average coordination scores, and channel topic-fixed effects; and a set of video-level predictors: video topic-fixed effects and a set of binary variables indicating whether the video contains hateful content, spreads misinformation or conspiracies, or was sent from a US location. Coefficient tables for these models are available in Appendix E.

In Figure 4 we focus on showing, from these different models, average marginal effect for our key covariates of interest: conservatism at the channel and video level. The left panel shows effects from models predicting channel suspensions. The other panels show effects from models predicting video suspensions: the middle one from models including all videos, and the right panel from models including only videos predicted to be about US politics. The most relevant distinction is that the orange effects are from models with no controls, only channel and/or video ideology as predictors; and the blue ones are from models with the controls mentioned above. The average marginal effects are for a one standard deviation increase in channel conservatism, and for conservative videos (v. liberal ones).

The models predicting channel and video suspension only as a function of the ideological leaning of the channel/video, show that conservative channels and videos are more likely to be suspended. The orange estimate on the left panel, indicates that a one standard deviation increase in channel conservatism increases the likelihood of channel suspension by 0.46 percentage points. The orange estimates in the middle and right panels show that the same

Figure 4: **A.** Marginal means from 6 logistic regressions predicting channel and video suspensions. **B.** Proportion of videos from channels from each ideological decile that are hateful (top) and spread misinformation (bottom). **C.** Most predictive words for hateful and misinformation/conspiracy videos.



standard deviation increase in channel conservatism correlates with a 0.91 and 0.46 percentage point increase in the likelihood of all videos, and political videos (respectively), being suspended. This smaller marginal effect of channel conservatism on the suspension of political videos is likely explained by Model 3 also accounting for the ideological leaning of the video. The bottom orange estimate in the right panel shows the strongest effect: political videos with a conservative leaning (v. liberal) are 1.69 percentage points more likely to be suspended (95% Confidence Interval: 0.57pp - 2.82pp). These findings are consistent with the narrative of some Republican elites that mainstream platforms such as YouTube have a pro-liberal bias when moderating content from politically-relevant sources.

However, in statistical models controlling for other theoretically relevant confounders, we find little-to-no evidence supporting the claim of asymmetric political moderation. Channel ideology is no longer predictive of channel suspension (blue estimate in the left panel, Model 4: 0.26pp, CI: $-0.03\text{pp} - 0.55\text{pp}$). The effect on video suspensions is close to 0 when all videos are included (blue estimate in the middle panel, Model 5: 0.47pp, CI: $0.17\text{pp} - 0.77\text{pp}$) and indistinguishable from 0 when only looking at political videos (top blue estimate in the

right panel, Model 6: 0.32pp, CI: $-0.05\text{pp} - 0.68\text{pp}$). Similarly, in this last model we see the difference in the likelihood of conservative (v. liberal) channels to be suspended to have substantially shrunk and to barely be distinguishable from 0 (0.7pp, CI: $0.09\text{pp} - 1.3\text{pp}$). As a robustness test, we estimate again this last model predicting the suspension of political videos, using designed-based supervised learning (DSL) (19) to ensure that the results are not a function of any systematic error or bias in our VLM-predicted variables. This additional model yields the same exact results (see Model 6 DSL in Appendix E).

Noticeably, in the models with controls, other confounders emerge as key predictors of suspension. For channels: number of days since the creation of the account and historical number of views are negative predictors, and misinformation and coordination are strong positive predictors of suspension. Similarly, for videos: number of days since the creation of the account and historical number of views are also negative predictors (as well as the account belonging to a politician or a mainstream media organisation); and likely indicators of platform violations, such as account coordination and posting hateful videos, are strong positive predictors of suspension.

Discussion

The moderation and curation of political speech on social media platforms is a topic of increasing political contention (20; 21; 22; 23). As additional examples, there are growing concerns about TikTok potentially amplifying posts in favor of the Chinese government (or censoring/shadowbanning those against it) (24), and about the new TikTok USDS company curating content in line with the current US administration (25). Yet data access and methodological challenges prevent academics, society, and policymakers from empirically assessing how platforms actually moderate and curate political content (26). This study uses state-of-the-art computational methods, such as Vision Language Models (VLMs), to collect and analyze

large amounts of politically-relevant data from YouTube for more than a year leading up to the 2024 US election, crucially contributing novel and detailed analyses to the ongoing debate on platform moderation biases. Contrary to the narrative of prominent Republican figures during the Biden Administration, this study reveals that ideological differences in the suspension of YouTube accounts and content during this period can be mostly explained by differences in compliance with the platform’s terms of service, rather than biased ideological decisions by the platform. In statistical models predicting channel and video suspensions as a function of the ideological leaning of the channel and videos, plus a range of theoretically relevant confounders, we observe little-to-no difference in the likelihood of content from different ideologies to be suspended. However, in descriptive analyses we do find that conservative channels, particularly those on the extreme, produced hateful and conspiracy videos at substantially higher rates; types of platform violations that in the multivariate analyses emerged as relevant predictors of suspension.

This study suffers from three main limitations. The first is the observational nature of the data. On the one hand, this is one of the key strengths of the study: to analyze at scale how a relevant platform such as YouTube *actually* moderates political content. However, as indicated above, the ideological leaning of the accounts and videos can be correlated with other relevant predictors of suspension. Isolating a potential correlation (or lack thereof) between suspensions and ideology depends on our ability to account for all relevant confounders. Building on past research (11; 27; 9; 8; 14), we have identified and measured many potential confounders, such as: whether videos are about politics, the specific policy areas discussed (22 policy areas), the topical focus of the channel (23 topic categories), account activity and engagement (number of videos, subscribers, and cumulative likes), number of days since the creation of the account (often found to be the strongest predictor of bot and malicious behavior (28; 9)), coordinated behavior (9), and whether the videos contain hateful content and spread conspiracies or misinformation. We believe that this long list of controls helped to account for

the most relevant confounders. A second, related limitation is that data-collection restrictions faced by independent researchers mean that our data likely miss some suspended videos. We are able to refresh the data for a given channel only about once a week, so any video that was posted and removed within that window would not be captured. However, because the timing of each channel’s update is not fixed in advance, this missingness should occur at random and is unlikely to have strong downstream effects. A final relevant limitation is that the analysis and conclusions presented here only apply to a particular platform and time, as happens with any cutting edge research in this area (29; 30; 8). Yet, the findings add robustness to, and align with, several previous studies looking at mainstream Western platforms, which also show: a higher volume of misinformation and conspiracies among (extreme) conservative social media accounts (31; 32), how a similar supply does not seem to exist on the left (17), and how this asymmetry can account for most ideological differences in platform moderation (8).

Future research should pay closer attention to other less visible forms of content moderation and curation, and whether they are subject to any kind of political bias. Platforms increasingly regulate violating content and accounts through more soft measures, such as down-ranking particular accounts or content, hiding them from search results, or disabling commenting or sharing features (27; 2). Research should also examine growing concerns among policymakers about whether TikTok, the first mainstream Western platform of non-Western origin, amplifies particular geopolitical narratives at the expense of others, an argument that was invoked, for example, to force the transfer of TikTok’s operations in the US (33).

A more normative discussion, is whether social media platforms should moderate content at all. Some political elites, particularly in the (far) right, increasingly advocate for non-moderated platforms such as Parler, Odysee, and Rumble (34; 23). However, at least in the aggregate, there is strong public support for some basic moderation in many countries, including in the US (21; 35; 36). Future research should explore the political and democratic consequences of non-moderated platforms, as well as the relationship between moderated and

non-moderated platforms; and whether moderation actions by moderated platforms push extreme content creators and users to more niche platforms (37; 38; 39), and the potential negative consequences this can have for society.

Finally, in order to produce more regular and cross-platform knowledge about how platforms moderate and curate (political) information, better academic access to platform data is essential. In the absence of improved access, policymakers and society rely on platform transparency reports, which are rare, vague, and unverifiable. Without clear independent analyses like the one presented here, unfounded claims can easily proliferate, contributing to fueling an already-polarized society, and leaving policymakers in the dark.

Methods

Sample

We identify and monitor a large and diverse pool of salient YouTube channels posting about US politics. We used a snowballing technique for building this sample (9; 40). The starting point was an extensive seed list of YouTube channels of media organizations (such as New York Times and Fox News. $N = 105$) and of US politicians (members of the 118th Congress, President Biden, and President Trump. $N = 184$). Then, in April 2023 we identified a set of politically-engaged users who commented or replied to videos from these elite channels ($N = 26,353$ users), and collected the full list of channels to which they were subscribed ($N = 7,334,005$ channels – 1,624,136 of them unique). We narrowed the list to focus on channels at least followed by 10 of our 26,353 users ($N = 92,653$), in order to make data collection more feasible, but also to focus on those channels that had a more substantive reach among these politically-interested users. Finally, we trained a BERT language model that used channel descriptions and tags to predict whether they posted about politics (95% Accuracy; 81%

Precision; 80% Recall). To account for the wide range of data sources from which people can get political information (41), we trained the model to identify channels as political even if only some of the content was of political relevance (e.g. entertainment channels, such as late night shows, that sometimes make political commentaries). Through this process we identified 20,054 politically-relevant channels, and added them to the original seed list of channels from political elites, for a total of 20,343. To account for new relevant channels that could have emerged since we first started collecting data, in February 2024 we repeated the same process, and topped up our sample with a few newly-discovered channels, for a total of 20,388 channels.

Data Collection

From July 1st 2023, through November 6th 2024 (the day after the election), we monitor these channels using the YouTube API. We developed a set of computer scripts to collect, on a rolling basis (about once a week per channel), the following information: (a) whether a channel is still active, (b) if not, the reason provided by YouTube for why it has been removed, (c) if active, all videos (metadata, transcript, and unique video frames) posted since the last time we checked (or a random sample of 100 videos if they had posted >100 since the last check), (d) whether the previously collected videos are still active, and (e) if not, the reason provided by YouTube for why they had been removed. The final binary suspension variables used in the analysis are based on a last suspension check we performed, both for channels and videos, at the end of December 2024. We consider channels and videos that had been temporary removed at any point, but that were active at the end of data collection, as non-suspended for the purpose of our analysis. Channels and videos can become unavailable because the user or the platform has taken them down. In the binary suspension measure we want to capture instances where the platform has taken down the content. For that purpose, we code a video as suspended (= 1) if the platform returns one of the following reasons: Channel terminated, Copyright Infringement, Not available, Removed due to 3rd Party Privacy Claim, Violation of

Community Guidelines, Violation of Harassing & Offensive Policy, Violation of Hate Speech Policy, Violation of Nudity & Sexual Policy, Violation of Terms of Service, and Violation of Violent Content Policy; and code videos as not suspended ($= 0$) if the platform returns one of these other reasons: Active (temporarily suspensions), Age protected, Channel terminated by the user, Private, Removed by the user.

Channel Ideology

We use Correspondence Analysis (CA) to estimate the ideology of the channels in our sample in a left-right continuum, adapting a widely validated and used method for estimating the ideology of Twitter accounts (15; 42; 9). We first built a network graph with information about which of the elite seed channels each of our politically-interested users was subscribed to – and then restricted the graph to users who followed at least 3 elite channels, resulting into a (8,548 users)x(289 elites) matrix. We then fit a CA model to the data using the `CA()` function of the `tweetscores` package in R. Finally, we projected the fitted model into a (8,548)x(20,348) matrix with information about which other channels from our sample these users were subscribed to, obtaining an ideology score for these 20,348 channels. In Appendix A, we explore the ideological scores obtained for the media and political elites in our sample. The ideological classification of media accounts has great face validity (Figure A1.A), and the ideological score for members of Congress correlates highly with other ideological measures such as their DW-NOMINATE score (A1.B and A1.C).

Channel Coordination

We measure the similarity of the content posted by a given channel and the content posted by the other channels in the dataset, as a likely marker of coordination. To do so, we first use a multilingual sentence transformer, BGE-M3, to extract a single transcript embedding for each video in our dataset. Due to context length restrictions of the model, we only use up

to the first 1,000 words from each transcript. Next, we average the transcript embeddings for all videos from a given channel, obtaining a channel-level embedding reflecting the content posted by that channel during the period of analysis. Next, we calculate the cosine similarity between all potential pairs of channel embeddings. Finally, we use these cosine similarities to obtain two measures of coordination: the mean similarity between the embedding of a given channel and the embeddings for all the other channels in the data; as well the maximum similarity with the embedding of another channel. In our multivariate models we include both coordination measures, as they could potentially be used by the platform for identifying coordinated accounts.

Automatic content classifiers

We trained five open-source Large Vision Language Models (VLMs) to automatically classify the content of *all English* videos tracked for one year leading to election day (Nov. 2023 - Nov. 2024), for five theoretical concepts of interest: (1) whether it contained **hateful** content (binary), (2) whether it spread **misinformation** (binary), and (3) whether it was about US **politics** (binary). In addition, we also classified the *subset of political videos* for (4) their **ideology** (categorical: Neutral, Liberal, Moderate, and Conservative) and (5) their particular political **topic** (categorical: *Comparative Agendas Project* classification). First, we manually annotated 3,884 videos for these five variables. All annotated videos are in English, and this is why in the main paper we restrict parts of the analysis only to videos in English, because those are the ones we confidently know we can measure with the measurement error reported in Appendix D. In selecting the videos to annotate, we over-sampled on suspended videos and also on videos predicted to be toxic by an off-the-shelf model (**detoxify**), aiming to inflate the number of true positives for the hateful and misinformation in the annotated dataset, and to fine-tune more accurate classifiers. After experimenting with several open-source models (SVM, BERT, Llama2, Llama3, Idefics2, and Idefics3), we settled on **Idefics3-8B-Llama3**,

an open-source VLM from HuggingFace that uses Llama3 as the text backbone. Both for fine-tuning and for inference, we use the first 1,000 tokens of the transcripts, and the first deduplicated frames (after having extracted 1 frame/second), from each video (the exact number of frames is an hyperparameter that varies by classifier, as indicated in Appendix D). In Appendix B, we show the distribution of the annotated data. Apart from the *US Politics* variable, despite the efforts in inflating rare categories for annotation, we still see most variables to be quite unbalanced. In Appendix D, we show the out-of-sample performance of each VLM, based on two different folds/random-seeds, and an untouched test set. We see model accuracy to be very high for all five VLMs. Precision and recall are very high for the *US Politics* and *Misinformation* model, but lower for the *Hateful* tasks. However, although only 48% of the predicted hateful messages are indeed hateful, the model does a really good job, particularly given that this a rare and complex task, to identify videos that are indeed hateful (72% recall). Finally, the performance for the multiclass classifiers (*Topic* and *Ideology*) is very high ($\sim 75\%$), particularly given the number of classes and the complexity of the task. In Appendix F, we show the prompts used for fine-tuning the model. In Appendix Appendix C we show Inter Rater Reliability (IRR) statistics based on two coders that annotated the same subset videos ($N = 310$). Despite some variables being highly unbalanced, Cohen’s Kappa for all is above 0.6 (often used as the threshold for *substantial* agreement), and in most cases between 0.7-0.8. The percentage agreement is above 70% for all variables, with agreement for most variables being between 80-100%.

Funding Declaration

This research has been possible thanks to a VENI grant from NWO (VI.Veni.211R.052, PI: Andreu Casas), and from two Small Computing Grants (EINF8284 and EINF53) and one Large Computing Grant (2024.044) from NWO, and a Future Leaders Fellowship from UKRI

(UKRI2362).

Author Contribution

A.C. conceived the study, collected and analyzed the data, developed and validated the computational models, and wrote the manuscript.

Data Availability

Code and data for replicating the analytical tables and figures in the manuscript will be made available on a public GitHub repository and on the Harvard Dataverse upon publication; and to reviewers and editors during the review process if requested.

Competing Interest

Authors declare no competing interests.

References


- [1] J. M. Balkin, Free speech in the algorithmic society: big data, private governance, and new school speech regulation, *UCDL Rev.* **51**, 1149–1210 (2017).
- [2] R. Gorwa, *The Politics of Platform Regulation: How Governments Shape Online Content Moderation* (Cambridge University Press, 2024).
- [3] C. S. Aubin, J. Liedke, Social media and news fact sheet, *Pew Research Center* pp. <https://www.pewresearch.org/journalism/fact--sheet/social--media--and--news--fact--sheet/> (2024).
- [4] T. Cruz, Sen. cruz: The pattern of political bias from YouTube and Google is massive, Press Release, U.S. Senator Ted Cruz of Texas (2019).
- [5] J. Jordan, Letter to Mark Zuckerberg, CEO, Meta Platforms: Big Tech censorship investigation, Letter, U.S. House Committee on the Judiciary, Ranking Member (2022).
- [6] Trump, Donald J., Restoring freedom of speech and ending federal censorship, Executive Order, The White House (2025).
- [7] E. A. Vogels, A. Perrin, M. Anderson, Most Americans think social media sites censor political viewpoints, Pew Research Center (2020).
- [8] M. Mosleh, Q. Yang, T. Zaman, G. Pennycook, D. G. Rand, Differences in misinformation sharing can lead to politically asymmetric sanctions, *Nature* **634**, 609–616 (2024).
- [9] A. Casas, The geopolitics of deplatforming: A study of suspensions of politically-interested iranian accounts on twitter, *Political Communication* **0**, 1-22 (2024).
- [10] Y. Yang, T. Davis, M. Hindman, Visual misinformation on facebook, *Journal of Communication* **73**, 316-328 (2023).
- [11] M. Bastos, This account doesnt exist: Tweet decay and the politics of deletion in the brexit debate, *American Behavioral Scientist* **65**, 757–773 (2021).
- [12] G. King, J. Pan, M. E. Roberts, How censorship in china allows government criticism but silences collective expression, *American political science Review* **107**, 326–343 (2013).
- [13] J. Earl, T. V. Maher, J. Pan, The digital repression of social movements, protest, and activism: A synthetic review, *Science Advances* **8**, eabl8198 (2022).
- [14] S. Majo-Vazquez, M. Congosto, T. Nicholls, R. K. Nielsen, The role of suspended accounts in political discussion on social media: Analysis of the 2017 french, uk and german elections, *Social Media + Society* **7**, 20563051211027202 (2021).

- [15] P. Barberá, Birds of the same feather tweet together: Bayesian ideal point estimation using twitter data, *Political analysis* **23**, 76–91 (2015).
- [16] G. Eady, R. Bonneau, J. A. Tucker, J. Nagler, News sharing on social media: Mapping the ideology of news media content, citizens, and politicians, *OSF Preprint*: osf.io/ch8gj (2020).
- [17] Y. Benkler, R. Faris, H. Roberts, *Network propaganda: Manipulation, disinformation, and radicalization in American politics* (Oxford University Press, 2018).
- [18] K. Ahuja, *et al.*, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, H. Bouamor, J. Pino, K. Bali, eds. (Association for Computational Linguistics, Singapore, 2023), pp. 4232–4267.
- [19] N. Egami, M. Hinck, B. Stewart, H. Wei, Using imperfect surrogates for downstream inference: Design-based supervised learning for social science applications of large language models, *Advances in Neural Information Processing Systems* **36** (2023).
- [20] A. A. Siegel, *Social Media and Democracy: The State of the Field, Prospects for Reform*, N. Persily, J. A. Tucker, eds. (Cambridge University Press, 2020), pp. 56–88.
- [21] A. Kozyreva, *et al.*, Resolving content moderation dilemmas between free speech and harmful misinformation, *Proceedings of the National Academy of Sciences* **120**, e2210666120 (2023).
- [22] F. Pradel, J. Zilinsky, S. Kosmidis, Y. Theocharis, Toxic speech and limited demand for content moderation on social media, *American Political Science Review* **118**, 1895–1912 (2024).
- [23] R. E. Appel, J. Pan, M. E. Roberts, Partisan conflict over content moderation is more than disagreement about facts, *Science Advances* **9**, eadg6799 (2023).
- [24] D. Leonhardt, Tiktok’s pro-china tilt, *The New York Times* (2024). Opinion.
- [25] J. Hendrix, Who will own tiktok in the us and why it matters for democracy, *Tech Policy Press* (2025). Analysis.
- [26] A. Casas, G. Dagher, B. O’Loughlin, Academic access to social media data for the study of political online safety, *Tech. rep.*, Center for Open Science (2025). Preprint.
- [27] K. Jaidka, S. Mukerjee, Y. Lelkes, Silenced on social media: the gatekeeping functions of shadowbans in the American Twitterverse, *Journal of Communication* **73**, 163–178 (2023).
- [28] M. T. Bastos, D. Mercea, The brexit botnet and user-generated hyperpartisan news, *Social Science Computer Review* **37**, 38–54 (2019).

- [29] A. M. Guess, *et al.*, How do social media feed algorithms affect attitudes and behavior in an election campaign?, *Science* **381**, 398–404 (2023).
- [30] S. González-Bailón, *et al.*, Asymmetric ideological segregation in exposure to political news on facebook, *Science* **381**, 392–398 (2023).
- [31] A. Guess, J. Nagler, J. Tucker, Less than you think: Prevalence and predictors of fake news dissemination on facebook, *Science advances* **5**, eaau4586 (2019).
- [32] N. Grinberg, K. Joseph, L. Friedland, B. Swire-Thompson, D. Lazer, Fake news on twitter during the 2016 u.s. presidential election, *Science* **363**, 374–378 (2019).
- [33] A. W. Palmer, How tiktok became a diplomatic crisis, *The New York Times Magazine* (2022).
- [34] K. Breuninger, D. Mangan, Trump vows ‘free speech’ reform of government, universities, media, tech firms if elected in 2024, *CNBC* (2022).
- [35] Y. Theocharis, S. Kosmidis, J. Zilinsky, F. Quint, F. Pradel, Content warning: Public attitudes on content moderation and freedom of expression (2025). Report, TUM Think Tank.
- [36] S. Munzert, R. Traunmüller, P. Barber, A. Guess, J. Yang, Citizen preferences for online hate speech regulation, *PNAS Nexus* **4**, pgaf032 (2025).
- [37] C. Buntain, R. Bonneau, J. Nagler, J. A. Tucker, Youtube recommendations and effects on sharing across online social platforms, *Proceedings of the ACM on Human-Computer Interaction* **5**, 11:1–11:26 (2021).
- [38] B. D. Horne, M. Bowman, M. Z. Trujillo, M. Gruppi, C. Buntain, Supply and demand on alt-tech social media: A case study of bitchute, *Computational Communication Research* **7**, 1 (2025).
- [39] C. Buntain, M. Snegovaya, Post-january 6 deplatforming shows long-term effects on ideological polarization among twitter users, *PNAS Nexus* **4**, pgaf333 (2025).
- [40] P. Barberá, *et al.*, Who leads? who follows? measuring issue attention and agenda setting by legislators and the mass public using social media data, *American Political Science Review* **113**, 883–901 (2019).
- [41] M. Boukes, H. G. Boomgaarden, M. Moorman, C. H. de Vreese, At odds: Laughing and thinking? the appreciation, processing, and persuasiveness of political satire, *Journal of Communication* **65**, 721–744 (2015).
- [42] P. Barbera, J. T. Jost, J. Nagler, J. A. Tucker, R. Bonneau, Tweeting from left to right: Is online political communication more than an echo chamber?, *Psychological Science* **26**, 1531–1542 (2015). PMID: 26297377.

Appendix A Estimating the ideology of YouTube channels.

Figure A1: Validation of the CA-based ideology scores



Appendix B Data annotations for training five VLMs.

Table B1: Distribution of the data annotated for training the Large Vision Language Models

Task	Description	Values	N	%
US Politics	Whether the video is about, or relevant to, US politics	0	1,935	49.8%
		1	1,945	50.2%
		<i>N</i>	3,880	100.0%
Hateful	Whether the video contains hateful language/behavior	0	3,431	88.4%
		1	449	11.6%
		<i>N</i>	3,880	100%
Ideology	The ideological leaning of the video	Neutral	238	21.7%
		Conservative	476	41.9%
		Moderate	176	15.5%
		Liberal	247	20.9%
		<i>N</i>	1,137	100%
Misinfo	Whether the video contains misinformation/conspiracies	0	3,672	94.5%
		1	208	5.5%
		<i>N</i>	3,880	100%
Topic	Main <i>Comparative Agendas Project</i> topic discussed in the video	Economy	51	2.7%
		Civil Rights	501	26.2%
		Healthcare	56	2.9%
		Agriculture	2	0.1%
		Labor	30	1.6%
		Education	23	1.2%
		Environment	27	1.4%
		Energy	19	1.0%
		Immigration	43	2.3%
		Transportation	13	0.7%
		Law & Crime	121	6.3%
		Social Welfare	22	1.2%
		Housing	35	1.8%
		Commerce	54	2.8%
		Defense	76	4%
		Technology	80	4.2%
		Foreign Trade	9	0.5%
		Intl. Affairs	253	13.2%
		Gov. Operations	428	22.4%
		Public Lands	9	0.5%
		Gun Control	59	3.1%
		<i>N</i>	1,911	100%

Appendix C Inter rater reliability for five VLMs.

Table C1: Inter rater reliability for five annotation tasks, for two annotators who annotated a common subset ($N = 310$) of the messages reported in Appendix B

Variable	Type	Cohen's Kappa	% Agreement
US Politics	Binary (0,1)	0.76	88.71
Hateful	Binary (0,1)	0.71	95.81
Misinformation	Binary (0,1)	0.62	94.19
Topic	Categorical (23 categories)	0.66	70.00
Ideology	Ordered Categorical (4 categories)	0.69 (Weighted = 0.8)	78.49

Appendix D Out-of-sample VLM performance.

Table D1: Two-fold validation based on a completely untouched test set, for five VLMs fine-tuned for the study (Idefics3-8B-Llama3).

Variable	% Modal Class	% True Pos.	Accuracy	Precision	Recall	F1-Score	#Images	#Epochs
US Politics	50.20	50.20	0.96	0.95	0.97	0.96	4	3
Hateful	88.40	11.60	0.88	0.48	0.72	0.58	5	10
Misinformation	94.50	5.50	0.99	0.95	0.89	0.92	3	1
Topic	26.20	-	0.77	-	-	-	8	2
Ideology	41.90		0.74	-	-	-	1	3

Appendix E Coefficient Table for logistic regressions predicting suspensions.

Table E1: Logistic regressions predicting channel suspension. Only channels that posted in English between November 2023, and election day in November 2024. **Note:** * indicates effect is statistically significant at conventional < 0.05 level.

	Model 1	Model 2
	Channel	Channel
	(no controls)	(w. controls)
(Intercept)	-4.3634 (0.1003)*	-6.007 (3.8796)
(Channel) Ideology: Conservatism	0.3495 (0.0953)*	0.2289 (0.1292)
(Channel) Days since creation (log)		-0.5609 (0.1924)*
(Channel) Total num. of subscribers (log)		0.101 (0.0902)
(Channel) Total num. of videos (log)		0.0293 (0.1076)
(Channel) Total num. of views (log)		-0.1874 (0.0616)*
(Channel) Num. of videos (log)		-0.1478 (0.1598)
(Channel) Num. of political videos (log)		0.1889 (0.1746)
(Channel) Num. of hateful videos (log)		0.0622 (0.126)
(Channel) Num. of misinformation videos (log)		0.3607 (0.1788)*
(Channel) Avg. coordination score		-4.8016 (2.6143)
(Channel) Max. coordination score		11.6248 (4.1918)*
(Channel) Topic-fixed effects	No	Yes
N	8414	8414
LogLik	-588.41	-441.25
AIC	1180.83	994.51

Table E2: Logistic regressions predicting video suspension, with standard errors clustered by channel. Only English videos posted by non-suspended channels between November 2023, and election day in November 2024. **Note:** * indicates effect is statistically significant at conventional < 0.05 level.

	Model 3	Model 4	Model 5	Model 6	Model 6 (DSL)
	Video (no controls) (all videos)	Video (w. controls) (all videos)	Video (no controls) (political)	Video (w. controls) (political)	Video (w. controls) (political)
(Intercept)	-4.1423 (0.0866)*	-9.4964 (4.1676)*	-4.4787 (0.1388)*	-10.4302 (5.2487)*	-10.0870 (5.2525)
(Channel) Ideology: Conservatism	0.5509 (0.0828)*	0.2916 (0.0936)*	0.2841 (0.135)*	0.2098 (0.1231)	0.2060 (0.1230)
(Video) Conservative (v. Liberal)			0.9337 (0.2288)*	0.4731 (0.1922)*	0.5005 (0.1925)*
(Video) Moderate (v. Liberal)			0.4721 (0.1573)*	0.4031 (0.1404)*	0.4377 (0.1416)*
(Video) Neutral (v. Liberal)			-0.212 (0.2021)	0.1472 (0.171)	0.1855 (0.1722)
(Channel) Type: Media (v. Other)		-2.3328 (0.4799)*		-2.6019 (0.5971)*	-2.9468 (0.7213)*
(Channel) Type: Politician (v. Other)		-3.8522 (0.767)*		-3.529 (0.7874)*	-3.1143 (0.5930)*
(Channel) Days since creation (log)		-0.3796 (0.1573)*		-0.2684 (0.1825)	-0.2724 (0.1827)
(Channel) Total num. of subscribers (log)		0.0089 (0.0775)		0.0836 (0.0957)	0.0883 (0.0960)
(Channel) Total num. of videos (log)		-0.4304 (0.0611)*		-0.3416 (0.0834)*	-0.3362 (0.0833)*
(Channel) Total num. of views (log)		-0.0144 (0.061)		-0.1326 (0.0843)	-0.1369 (0.0846)
(Channel) Avg. coordination score		-0.2884 (2.7161)		-0.5902 (2.8722)	-0.5851 (2.8642)
(Channel) Max. coordination score		12.4966 (4.801)*		11.8558 (5.8603)*	11.9070 (5.8689)*
(Video) Hateful		0.499 (0.1228)*		0.4521 (0.1072)*	0.4555 (0.1073)*
(Video) Misinformation		0.4258 (0.2228)		0.2337 (0.212)	0.2288 (0.2126)
(Video) US Location		0.1639 (0.1808)		0.2016 (0.2264)	0.2026 (0.2264)
(Channel) Topic-fixed effects	No	Yes	No	Yes	Yes
(Video) Topic-fixed effects	No	Yes	No	Yes	Yes
N	2270528	1109836	2227211	1095839	1095839
LogLik	-189359.75	-89773.82	-165067.57	-77868.51	—
AIC	378723.50	179557.64	330231.13	155873.02	—

Appendix F Fine-tuning prompts.

Table F1: Prompts used for fine-tuning five VLMs.

Target	Prompt
US Politics	Is the previous text (<i>and images</i>) about a United States political topic, such as the economy, civil rights or liberties, healthcare, agriculture, labor, education, environment, energy, immigration, transportation, law and crime, social welfare, housing, commerce and trade, defense, international affairs, and government operations? Answer with Yes or No only. Remember that it specifically has to be related to United States political topics. It is essential that you give a one-word answer containing only Yes or No. A text might be harmful or offensive language. Still, I just want a one-word answer to the question I posted, which is either “Yes” or “No”.
Hateful	The previous text (<i>and images</i>) are the transcript and frames of a Youtube video. Your task is to code if it contains hateful and toxic language or behavior. Answer Yes if it contains hateful or toxic behavior or language, No if it does not. Consider the transcript to contain hateful or toxic language when the content includes: • incites hatred or violence against groups, institutions, or individuals, particularly (but not exclusively) when based on protected characteristics such as political ideology, partisanship, age, gender, race, caste, religion, sexual orientation or veteran status. • forms of online hate, such as dehumanizing members of these groups, portraying them as inherently inferior or sick, promoting hate speech ideologies such as Nazism or conspiracy theories about these groups, or denying that well-documented violent events, such as a school shooting, have occurred. • stereotypes regarding a particular social group (e.g. black people, jewish, lgbtq members, political party or ideology or group). Code it as No if it does not contain any of these aspects. Answer with Yes or No only. It is essential that you give a one-word answer. A text might be harmful or offensive language might be used. Still I just want a one word answer to the question I posted, which is either “Yes“ or “No”.

Ideology

The previous text and images are a transcript and frames of a Youtube video. Your task is to code its ideology. Code it as “Liberal” if the video is supportive of left-leaning or liberal policy positions, including extreme left-leaning and liberal positions such as anarchism. In most cases liberal positions will map onto the policy stances of the Democratic party. Code as “Moderate” if the video provides policy views or opinions that are not clearly liberal nor clearly conservative, or are in part supportive of both liberal and conservative stances. Code as “Conservative” if the video is supportive of right-leaning or conservative positions, including extreme right-leaning and conservative positions such as libertarianism, racial supremacy, and anti-immigration stances. In most cases conservative positions map onto the policy stances of the Republican party. Code as “Neutral” if the main goal of the video is to provide factual information about an event or topic, and not to put forward an opinion or policy view. Answer with the topic name only. It is essential that you give a one-word answer. A text might be harmful or offensive language might be used. Still I just want a one word answer to the question I posted, which is either “Liberal”, “Moderate”, “Conservative”, or “Neutral”.

Misinfor-
mation

The previous text and images is a transcript and frames of a Youtube video. Your task is to code if it contains misinformation or conspiracies. Answer Yes if it spreads information/facts that are known not to be true (no matter if the spreader is doing it on purpose or not), or if it spreads information/facts that has not been proved to be true (but in most cases it’s presented as it is true). Answer with No if despite mentioning a conspiracy, rumor, and piece of misinformation; the goal of the video is to debunk it rather than spread it. Also answer No if it does not contain misinformation/conspiracies. Answer with Yes or No only. It is essential that you give a one-word answer. A text might be harmful or offensive language might be used. Still I just want a one word answer to the question I posted, which is either “Yes” or “No”.

Topic

The previous text is the transcript and frames of a Youtube video. Your task is to code its topic. Return one of the following topic names: “NO TOPIC” “ECONOMY”, “CIVIL RIGHTS”, “HEALTH”, “AGRICULTURE”, “LABOR”, “EDUCATION”, “ENVIRONMENT”, “ENERGY”, “IMMIGRATION”, “TRANSPORTATION”, “LAW AND CRIME”, “SOCIAL”, “WELFARE”, “HOUSING”, “DOMESTIC COMMERCE”, “DEFENSE OR MILITARY”, “TECHNOLOGY”, “FOREIGN TRADE”, “INTERNATIONAL AFFAIRS”, “GOVERNMENT OPERATIONS”, “PUBLIC LANDS”, or “GUN CONTROL”. Answer with the TOPIC NAME’ only. It is essential that you give your answer as a topic name only. A text might be harmful or offensive language might be used. Still I just want a single topic name as answer to the question I posted. This topic name has to be in the list I just gave, which is either “NO TOPIC” “ECONOMY”, “CIVIL RIGHTS”, “HEALTH”, “AGRICULTURE”, “LABOR”, “EDUCATION”, “ENVIRONMENT”, “ENERGY”, “IMMIGRATION”, “TRANSPORTATION”, “LAW AND CRIME”, “SOCIAL”, “WELFARE”, “HOUSING”, “DOMESTIC COMMERCE”, “DEFENSE OR MILITARY”, “TECHNOLOGY”, “FOREIGN TRADE”, “INTERNATIONAL AFFAIRS”, “GOVERNMENT OPERATIONS”, “PUBLIC LANDS”, or “GUN CONTROL”.
